

Introduction to Machine Learning

Lecture 12: Convolutional Neural Network (CNN)

Nov 11, 2019

Jie Wang

Machine Intelligence Research and Applications Lab

Department of Electronic Engineering and Information Science (EEIS)

<http://staff.ustc.edu.cn/~jwangx/>

jiewangx@ustc.edu.cn



Machine Intelligence Research and Applications Lab



Introduction

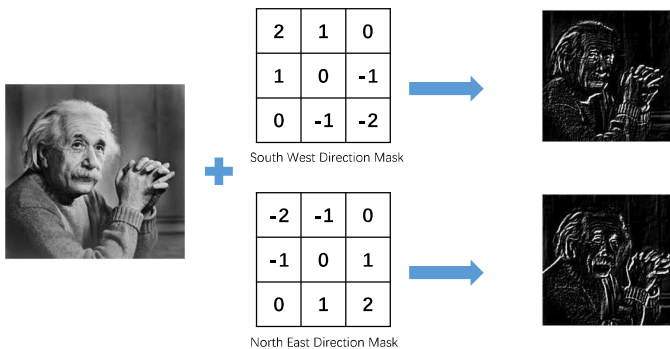
Features



$$\approx \underbrace{\text{img}_1 + \text{img}_2 + \text{img}_3 + \text{img}_4 + \text{img}_5 + \dots + \text{img}_n}_{\text{Features}} + \text{Relative location}$$

- How to extract discriminative features ?
 - Hand-crafted features: HOG, SIFT, SURF etc
 - Learned features: hidden states of DNN

Hand-crafted Features – Robinson Compass Mask

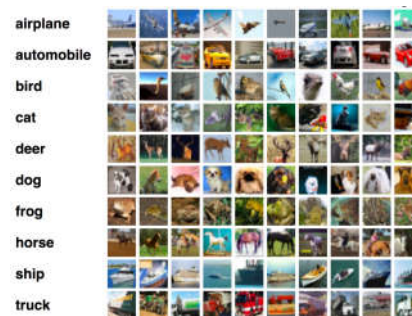


https://www.tutorialspoint.com/dip/robinson_compass_mask.htm

Contents

- Introduction
- Network Layers of CNN
- Learning a CNN
- Examples of CNN Architectures
- More Applications


Image Classification



- Identifying objects from various scenes is a easy task for human
- However, it is difficult for human to describe (precisely) how he/she can do it

<https://medium.com/@tifa2up/image-classification-using-deep-neural-networks-a-beginner-friendly-approach-using-tensorflow-94b0a090cc4>


Hand-crafted Features – Sobel Operator



https://www.tutorialspoint.com/dip/sobel_operator.htm

Hand-crafted Features – HOG

- Histograms of Oriented Gradients





<https://www.learnopencv.com/histogram-of-oriented-gradients/>

Dalal, Naveet, Triggs, et al. Histograms of Oriented Gradients for Human Detection. CVPR, 2005.

Hand-crafted Features – HOG


- Histograms of Oriented Gradients



<https://www.learnopencv.com/histogram-of-oriented-gradients/>
Dafal, Navneet, Triggs, et al. Histograms of Oriented Gradients for Human Detection. CVPR, 2005.

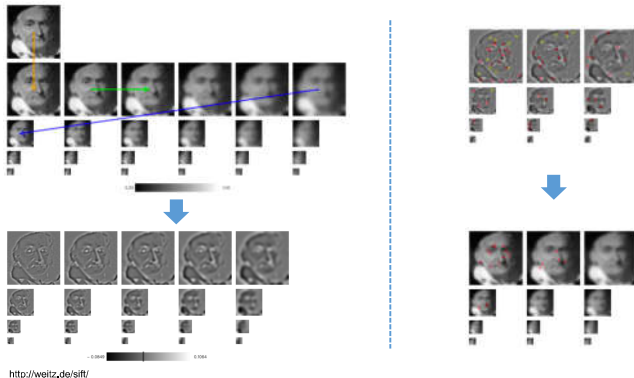
Hand-crafted Features – SIFT

- Scale-Invariant Feature Transform



Lowe D G. Distinctive Image Features from Scale-Invariant Keypoints[C]// International Journal of Computer Vision, 2004:91-110.
<http://www.cs.cmu.edu/~dst/SIFT/>

Hand-crafted Features – SIFT




Hand-crafted Features vs Learned Features



- Hand-crafted Features
 - Challenging to design
 - Require expert domain knowledge
 - Not flexible

- Learned Features
 - Automatically learned by machines
 - No need of expert domain knowledge
 - Data-driven

Learned Features



$$\approx \text{[Pattern 1]} + \text{[Pattern 2]} + \text{[Pattern 3]} + \text{[Pattern 4]} + \text{[Pattern 5]} + \dots + \text{Relative location}$$


- We can design a new neural network
 - It can detect different patterns
 - It can detect similar patterns in different regions
 - It can roughly preserve the spatial information

Detect different patterns



- Most patterns are much smaller in view of the whole image
 - A filter does not have to see the whole image to discover the pattern
 - One filter connects to small region with less parameters at a time



Detect the similar patterns in different regions



- The similar patterns appear in different regions
 - We can use one filter to detect similar patterns
 - One filter uses the same set of parameters for different regions




The spatial information



- For many vision tasks, the detected spatial information can be redundant

Roughly preserve the spatial information




- The relative location of patterns will be preserved after subsampling
 - Using subsampling to lessen the parameters

Network Layers of CNN


Convolutional Neural Networks (CNN)

- Try to construct a new neural network
 - One filter connects to small region with less parameters at a time
 - One filter uses the same set of parameters for different regions
 - Using subsampling to lessen the parameters
- Convolution
- Pooling (Subsampling)



How does CNN work


- A simple but important CNN – LeNet 5



- Given an input image, how to predict its label?

How does CNN work

- A simple but important CNN – LeNet 5



CNN – Convolution Layer

Input: Image

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6x6 image

Filters

1	1	1
-1	1	-1
1	-1	-1

Filter 1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 2

Those are parameters to be learned!

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6x6 image

Filter 1

1	1	1
-1	1	-1
1	-1	-1

One filter Connects to small region with less parameters at a time

$1 \times 1 + 1 \times 1 + 1 \times 1$
 $+ 0 \times (-1) + 0 \times 1 + 0 \times (-1)$
 $+ 0 \times 1 + 1 \times (-1) + 1 \times (-1)$

1

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6x6 image

Stride = 1

Filter 1

1	1	1
-1	1	-1
1	-1	-1

One filter uses the same set of parameters for different regions

1 2

-1

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

Stride = 3

Filter 1

1	1	1
-1	1	-1
1	-1	-1

➤ One filter uses the same set of parameters for different regions

1	5
3	

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

Stride = 1

Filter 1

1	1	1
-1	1	-1
1	-1	-1

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

Stride = 1

Filter 1

1	1	1
-1	1	-1
1	-1	-1

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

Stride = 1

Filter 1

1	1	1
-1	1	-1
1	-1	-1

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

CNN – Convolution Layer

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

Stride = 1

Filter 2

1	-1	-1
-1	1	-1
-1	-1	1

-1	-2	-4	-1
1	2	2	-2
-4	-1	-1	1
-1	-2	1	0

2×4×4 Images

Feature Maps

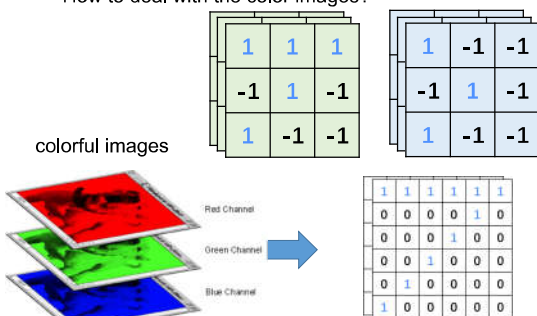
Ignore the bias and activation function!

CNN – Convolution Layer

1. How to deal with the color images?
2. How many parameters in the convolution layer?
3. What is the size of the feature maps?


CNN – Convolution Layer

- How to deal with the color images?



CNN – Convolution Layer

- How to deal with the colorful images?



CNN – Convolution Layer

- How many parameters in the convolution layer?

1	1	1
-1	1	-1
1	-1	-1

Filter 1

$$9 = 3 \times 3$$

1	-1	-1
-1	1	-1
-1	-1	1

Filter 2

$$9 = 3 \times 3$$

-1	-1	-1
-1	1	-1
1	-1	1

Filter N

$$9 = 3 \times 3$$

There are $9N$ parameters

The power of sharing weights!

CNN – Convolution Layer

- Padding

Stride = 2

1	1	1	1	1	1	0
0	0	0	0	0	1	0
0	1	1	1	0	0	0
0	0	1	0	0	0	0
0	1	0	0	0	0	0
1	0	0	0	0	0	0

6×6 image

1	1	1
-1	1	-1
1	-1	-1

- Two ways

1. Ignore the extra pixels

2. Fill with zero

Add P_H and P_W zeros

$$\text{Feature maps} = N \times \left(\frac{H - n + P_H}{s} + 1 \right) \times \left(\frac{W - n + P_W}{s} + 1 \right)$$

参考 cs231n 2017 lecture5, page52和
https://www.tensorflow.org/api_guides/python/nn#Notes_on_SAME_Convolution_Padding

CNN – Pooling Layer

- The output from the convolution layer can be huge

Input image: $1 \times 96 \times 96$

Filters: $400 \times 8 \times 8$

Stride: 1

Feature maps = $400 \times (96 - 8 + 1) \times (96 - 8 + 1)$

- Output of the convolution layer

$$3168400 = 400 \times (96 - 8 + 1) \times (96 - 8 + 1)$$

- Hard to train
- Overfitting

CNN – Pooling Layer

- Average Pooling

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

Convolved Feature

1	1	1
-1	1	-1
1	-1	-1

Filter 1

0	2.25
2	0.75

Pooled Feature

CNN – Convolution Layer

- What is the size of the feature maps?

1	1	1	1	1	1
0	0	0	0	1	0
0	1	1	1	0	0
0	0	1	0	0	0
0	1	0	0	0	0
1	0	0	0	0	0

6×6 image

2 filters
Filter size: 3×3
Stride:1

-1	-2	-4	-1
1	0	0	-2
-4	-1	-1	1
-1	-2	1	0

2×4×4 image

Input image: $C \times H \times W$

Filters: $N \times n \times n$

Stride: s


Feature maps=?

$$N \times \left(\frac{H - n}{s} + 1 \right) \times \left(\frac{W - n}{s} + 1 \right)$$

May not be divisible!

CNN – Pooling Layer

- LeNet 5



CNN – Pooling Layer

- Max Pooling

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

Convolved Feature

1	1	1
-1	1	-1
1	-1	-1

Filter 1

2	5
5	1

Pooled Feature

Using subsampling to lessen the parameters

CNN – Pooling Layer

- Overlapping Pooling

1	2	2	5
-1	-2	2	0
0	5	1	1
3	0	1	0

Convolved Feature

1	1	1
-1	1	-1
1	-1	-1


Filter 1

2	2	5
5	5	2
5	5	1

Pooled Feature


CNN - Convolution + Pooling

- LeNet 5



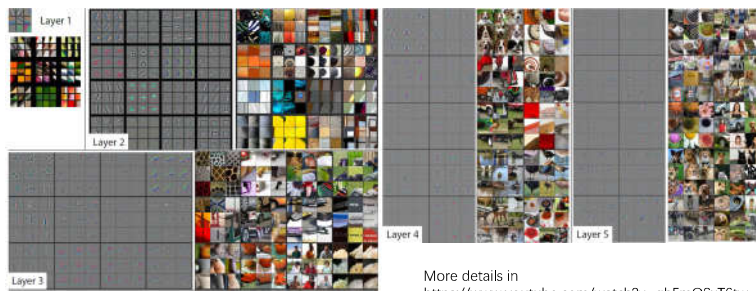
CNN - Convolution + Pooling

- LeNet 5



- Why more convolution and pooling layer?

Learned Features by CNN




More details in <https://www.youtube.com/watch?v=ghEmQSxT6tw>

- Deeper layers, more specific features


Zeiler M D, Fergus R. Visualizing and Understanding Convolutional Networks[J]. 2013, 8689:818-833.

CNN – Fully connected layer


- LeNet 5



CNN – Fully connected layer



CNN – Visualization



3D convolutional network visualization <http://scs.ryerson.ca/~aharley/vis/conv/>
A. W. Harley, "An Interactive Node-Link Visualization of Convolutional Neural Networks," in ISVC, pages 867-877, 2015

CNN – Hyperparameters

- Convolution layers
 - Number of filters
 - Size of filters
 - Stride
- Pooling layers
 - Window size
 - Window stride
- Fully connected layers
 - Number of layers
 - Number of neurons

CNN – Traditional Methods

- Comparison


	Traditional Methods	CNN
Filters	Manually design	Learn automatically
Layers	Few	Can be quite some
Features	Low level features	From low to high level features

CNN is more powerful!

Learning a CNN


Backpropagation

- A simple but important CNN – LeNet 5



- Back-propagation – Chain rule
- How to compute the gradients of convolution layers and pooling layers?

Backpropagation – Convolution Layer



- How to compute $\frac{\partial E}{\partial F_{11}}$?

Chain Rule

$$\frac{\partial E}{\partial F_{11}} = \left(\frac{\partial E}{\partial O} \right)^T \frac{\partial O}{\partial F_{11}}$$

$$\begin{aligned} \frac{\partial E}{\partial O} &= \nabla E(O_{11}, O_{12}, O_{21}, O_{22}) \\ &= \left(\frac{\partial E}{\partial O_{11}}, \frac{\partial E}{\partial O_{12}}, \frac{\partial E}{\partial O_{21}}, \frac{\partial E}{\partial O_{22}} \right)^T \end{aligned}$$

$$O = f(F, X) = (O_{11}, O_{12}, O_{21}, O_{22})^T \triangleq (f_1(F, X), f_2(F, X), f_3(F, X), f_4(F, X))^T$$

$$\frac{\partial O}{\partial F_{11}} = \left(\frac{\partial f_1}{\partial F_{11}}, \frac{\partial f_2}{\partial F_{11}}, \frac{\partial f_3}{\partial F_{11}}, \frac{\partial f_4}{\partial F_{11}} \right)^T$$

Backpropagation – Convolution Layer

$$O = f(F; X) \quad f_1(X) = O_{11} = F_{11}X_{11} + F_{12}X_{12} + F_{21}X_{21} + F_{22}X_{22}$$

$$f_2(X) = O_{12} = F_{11}X_{12} + F_{12}X_{13} + F_{21}X_{22} + F_{22}X_{23}$$

$$f_3(X) = O_{21} = F_{11}X_{21} + F_{12}X_{22} + F_{21}X_{31} + F_{22}X_{32}$$

$$f_4(X) = O_{22} = F_{11}X_{22} + F_{12}X_{23} + F_{21}X_{32} + F_{22}X_{33}$$

$$\begin{array}{cccc} \frac{\partial f_1}{\partial F_{11}} = X_{11} & \frac{\partial f_2}{\partial F_{11}} = X_{12} & \frac{\partial f_3}{\partial F_{11}} = X_{21} & \frac{\partial f_4}{\partial F_{11}} = X_{22} \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial f_1}{\partial F_{22}} = X_{22} & \frac{\partial f_2}{\partial F_{22}} = X_{23} & \frac{\partial f_3}{\partial F_{22}} = X_{32} & \frac{\partial f_4}{\partial F_{22}} = X_{33} \end{array}$$

Loss functions

- Mean squared error (MSE)


$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

- Cross entropy loss

$$\text{Loss} = - \sum_{i=1}^n Y_i \log p_i$$

- User defined loss

Backpropagation – Convolution Layer



- The output of convolution operation $O = f(F, X)$
- The loss E
- Assume that we have already computed $\frac{\partial E}{\partial O_{ij}}$, and of course all partial derivatives of latter layers

Backpropagation – Convolution Layer

$$\frac{\partial E}{\partial F_{11}} = \frac{\partial E}{\partial O_{11}} \frac{\partial f}{\partial F_{11}} + \frac{\partial E}{\partial O_{12}} \frac{\partial f}{\partial F_{11}} + \frac{\partial E}{\partial O_{21}} \frac{\partial f}{\partial F_{11}} + \frac{\partial E}{\partial O_{22}} \frac{\partial f}{\partial F_{11}}$$

$$\frac{\partial E}{\partial F_{12}} = \frac{\partial E}{\partial O_{11}} \frac{\partial f}{\partial F_{12}} + \frac{\partial E}{\partial O_{12}} \frac{\partial f}{\partial F_{12}} + \frac{\partial E}{\partial O_{21}} \frac{\partial f}{\partial F_{12}} + \frac{\partial E}{\partial O_{22}} \frac{\partial f}{\partial F_{12}}$$

$$\frac{\partial E}{\partial F_{21}} = \frac{\partial E}{\partial O_{11}} \frac{\partial f}{\partial F_{21}} + \frac{\partial E}{\partial O_{12}} \frac{\partial f}{\partial F_{21}} + \frac{\partial E}{\partial O_{21}} \frac{\partial f}{\partial F_{21}} + \frac{\partial E}{\partial O_{22}} \frac{\partial f}{\partial F_{21}}$$

$$\frac{\partial E}{\partial F_{22}} = \frac{\partial E}{\partial O_{11}} \frac{\partial f}{\partial F_{22}} + \frac{\partial E}{\partial O_{12}} \frac{\partial f}{\partial F_{22}} + \frac{\partial E}{\partial O_{21}} \frac{\partial f}{\partial F_{22}} + \frac{\partial E}{\partial O_{22}} \frac{\partial f}{\partial F_{22}}$$

Backpropagation – Convolution Layer

$$\frac{\partial E}{\partial F_{11}} = \frac{\partial E}{\partial O_{11}} X_{11} + \frac{\partial E}{\partial O_{12}} X_{12} + \frac{\partial E}{\partial O_{21}} X_{21} + \frac{\partial E}{\partial O_{22}} X_{22}$$

$$\frac{\partial E}{\partial F_{12}} = \frac{\partial E}{\partial O_{11}} X_{12} + \frac{\partial E}{\partial O_{12}} X_{13} + \frac{\partial E}{\partial O_{21}} X_{22} + \frac{\partial E}{\partial O_{22}} X_{23}$$

$$\frac{\partial E}{\partial F_{21}} = \frac{\partial E}{\partial O_{11}} X_{21} + \frac{\partial E}{\partial O_{12}} X_{22} + \frac{\partial E}{\partial O_{21}} X_{31} + \frac{\partial E}{\partial O_{22}} X_{32}$$

$$\frac{\partial E}{\partial F_{22}} = \frac{\partial E}{\partial O_{11}} X_{22} + \frac{\partial E}{\partial O_{12}} X_{23} + \frac{\partial E}{\partial O_{21}} X_{32} + \frac{\partial E}{\partial O_{22}} X_{33}$$

Backpropagation – Convolution Layer

X_{11}	X_{12}	X_{13}
X_{21}	X_{22}	X_{23}
X_{31}	X_{32}	X_{33}

+

$\frac{\partial E}{\partial O_{11}}$	$\frac{\partial E}{\partial O_{12}}$
$\frac{\partial E}{\partial O_{21}}$	$\frac{\partial E}{\partial O_{22}}$



$\frac{\partial E}{\partial F_{11}}$	$\frac{\partial E}{\partial F_{12}}$
$\frac{\partial E}{\partial F_{21}}$	$\frac{\partial E}{\partial F_{22}}$

$$\begin{aligned}\frac{\partial E}{\partial F_{11}} &= \frac{\partial E}{\partial O_{11}} X_{11} + \frac{\partial E}{\partial O_{12}} X_{12} + \frac{\partial E}{\partial O_{21}} X_{21} + \frac{\partial E}{\partial O_{22}} X_{22} \\ \frac{\partial E}{\partial F_{12}} &= \frac{\partial E}{\partial O_{11}} X_{12} + \frac{\partial E}{\partial O_{12}} X_{13} + \frac{\partial E}{\partial O_{21}} X_{22} + \frac{\partial E}{\partial O_{22}} X_{23} \\ \frac{\partial E}{\partial F_{21}} &= \frac{\partial E}{\partial O_{11}} X_{21} + \frac{\partial E}{\partial O_{12}} X_{22} + \frac{\partial E}{\partial O_{21}} X_{31} + \frac{\partial E}{\partial O_{22}} X_{32} \\ \frac{\partial E}{\partial F_{22}} &= \frac{\partial E}{\partial O_{11}} X_{22} + \frac{\partial E}{\partial O_{12}} X_{23} + \frac{\partial E}{\partial O_{21}} X_{32} + \frac{\partial E}{\partial O_{22}} X_{33}\end{aligned}$$

Backpropagation – Convolution Layer

$$\begin{aligned}\mathbf{O} &= f(\mathbf{F}; \mathbf{X}) \quad f_1(\mathbf{X}) = O_{11} = F_{11}X_{11} + F_{12}X_{12} + F_{21}X_{21} + F_{22}X_{22} \\ f_2(\mathbf{X}) &= O_{12} = F_{11}X_{12} + F_{12}X_{13} + F_{21}X_{22} + F_{22}X_{23} \\ f_3(\mathbf{X}) &= O_{21} = F_{11}X_{21} + F_{12}X_{22} + F_{21}X_{31} + F_{22}X_{32} \\ f_4(\mathbf{X}) &= O_{22} = F_{11}X_{22} + F_{12}X_{23} + F_{21}X_{32} + F_{22}X_{33}\end{aligned}$$

$$\begin{aligned}\frac{\partial E}{\partial X_{11}} &= \frac{\partial E}{\partial O_{11}} F_{11} + \frac{\partial E}{\partial O_{12}} 0 + \frac{\partial E}{\partial O_{21}} 0 + \frac{\partial E}{\partial O_{22}} 0 & \frac{\partial E}{\partial X_{12}} &= \frac{\partial E}{\partial O_{11}} F_{12} + \frac{\partial E}{\partial O_{12}} F_{11} + \frac{\partial E}{\partial O_{21}} 0 + \frac{\partial E}{\partial O_{22}} 0 & \frac{\partial E}{\partial X_{13}} &= \frac{\partial E}{\partial O_{11}} 0 + \frac{\partial E}{\partial O_{12}} F_{12} + \frac{\partial E}{\partial O_{21}} 0 + \frac{\partial E}{\partial O_{22}} 0 \\ \frac{\partial E}{\partial X_{21}} &= \frac{\partial E}{\partial O_{11}} F_{21} + \frac{\partial E}{\partial O_{12}} 0 + \frac{\partial E}{\partial O_{21}} F_{11} + \frac{\partial E}{\partial O_{22}} 0 & \frac{\partial E}{\partial X_{22}} &= \frac{\partial E}{\partial O_{11}} F_{22} + \frac{\partial E}{\partial O_{12}} F_{12} + \frac{\partial E}{\partial O_{21}} F_{11} + \frac{\partial E}{\partial O_{22}} F_{11} & \frac{\partial E}{\partial X_{23}} &= \frac{\partial E}{\partial O_{11}} 0 + \frac{\partial E}{\partial O_{12}} F_{22} + \frac{\partial E}{\partial O_{21}} 0 + \frac{\partial E}{\partial O_{22}} F_{11} \\ \frac{\partial E}{\partial X_{31}} &= \frac{\partial E}{\partial O_{11}} 0 + \frac{\partial E}{\partial O_{12}} 0 + \frac{\partial E}{\partial O_{21}} F_{21} + \frac{\partial E}{\partial O_{22}} 0 & \frac{\partial E}{\partial X_{32}} &= \frac{\partial E}{\partial O_{11}} 0 + \frac{\partial E}{\partial O_{12}} 0 + \frac{\partial E}{\partial O_{21}} F_{22} + \frac{\partial E}{\partial O_{22}} F_{21} & \frac{\partial E}{\partial X_{33}} &= \frac{\partial E}{\partial O_{11}} 0 + \frac{\partial E}{\partial O_{12}} 0 + \frac{\partial E}{\partial O_{21}} 0 + \frac{\partial E}{\partial O_{22}} F_{22}\end{aligned}$$

Backpropagation – Pooling Layer

- Max Pooling

X_{11}^i	X_{12}^i	X_{13}^i	X_{14}^i
X_{21}^i	X_{22}^i	X_{23}^i	X_{24}^i
X_{31}^i	X_{32}^i	X_{33}^i	X_{34}^i
X_{41}^i	X_{42}^i	X_{43}^i	X_{44}^i

4x4 Input



X_{11}^o	X_{23}^o
X_{42}^o	X_{33}^o

2x2 output

→ E

- Only need to compute the gradients of the inputs


$$\begin{aligned}\frac{\partial E}{\partial X_{11}^i} &= \frac{\partial E}{\partial X_{11}^o} & \frac{\partial E}{\partial X_{12}^i} &= 0 & \frac{\partial E}{\partial X_{21}^i} &= 0 & \frac{\partial E}{\partial X_{22}^i} &= 0 \\ \frac{\partial E}{\partial X_{13}^i} &= 0 & \frac{\partial E}{\partial X_{14}^i} &= 0 & \frac{\partial E}{\partial X_{23}^i} &= \frac{\partial E}{\partial X_{23}^o} & \frac{\partial E}{\partial X_{24}^i} &= 0 \\ & & & & & & & \vdots\end{aligned}$$

We assign the error to where it comes from

LeNet (1998)

- Gradient-based learning applied to document recognition

[Yann LeCun, Leon Bottou, Yoshua Bengio, and Patrick Haffner 1998]



LeNet-5

5 layers

Backpropagation – Convolution Layer

- How to compute $\frac{\partial E}{\partial X_{11}}$?

Chain Rule

$$\frac{\partial E}{\partial X_{11}} = \left(\frac{\partial E}{\partial \mathbf{O}} \right)^T \frac{\partial \mathbf{O}}{\partial X_{11}}$$

$$\begin{aligned}\frac{\partial E}{\partial \mathbf{O}} &= \nabla E(O_{11}, O_{12}, O_{21}, O_{22}) \\ &= \left(\frac{\partial E}{\partial O_{11}}, \frac{\partial E}{\partial O_{12}}, \frac{\partial E}{\partial O_{21}}, \frac{\partial E}{\partial O_{22}} \right)^T\end{aligned}$$

$$\mathbf{O} = f(\mathbf{F}, \mathbf{X}) = (O_{11}, O_{12}, O_{21}, O_{22})^T \triangleq (f_1(\mathbf{F}, \mathbf{X}), f_2(\mathbf{F}, \mathbf{X}), f_3(\mathbf{F}, \mathbf{X}), f_4(\mathbf{F}, \mathbf{X}))^T$$

$$\frac{\partial \mathbf{O}}{\partial X_{11}} = \left(\frac{\partial f_1}{\partial X_{11}}, \frac{\partial f_2}{\partial X_{11}}, \frac{\partial f_3}{\partial X_{11}}, \frac{\partial f_4}{\partial X_{11}} \right)^T$$

Backpropagation – Convolution Layer

0	0	0	0
0	$\frac{\partial E}{\partial O_{11}}$	$\frac{\partial E}{\partial O_{12}}$	0
0	$\frac{\partial E}{\partial O_{21}}$	$\frac{\partial E}{\partial O_{22}}$	0
0	0	0	0

+

F_{22}	F_{21}
F_{12}	F_{11}



$\frac{\partial E}{\partial X_{11}}$	$\frac{\partial E}{\partial X_{12}}$	$\frac{\partial E}{\partial X_{13}}$
$\frac{\partial E}{\partial X_{21}}$	$\frac{\partial E}{\partial X_{22}}$	$\frac{\partial E}{\partial X_{23}}$
$\frac{\partial E}{\partial X_{31}}$	$\frac{\partial E}{\partial X_{32}}$	$\frac{\partial E}{\partial X_{33}}$

Examples of CNN Architecture

ImageNet (Benchmark dataset)

- ImageNet

- About 1.5×10^7 images, 2.2×10^4 categories
- An image database organized according to the WordNet hierarchy



J Deng, W Dong, R Socher, J Li, K Li, L Fei-Fei. ImageNet: A large-scale hierarchical image. CVPR 2009.

The figure is divided into two main sections: Classification and Classification + Localization.







Classification:

- Steel drum:** An image of a person playing a steel drum. A list of objects includes Scale, T-shirt, Steel drum, Drumstick, and Mud turtle. A green checkmark indicates successful classification.
- Giant panda:** A list of objects includes Scale, T-shirt, Giant panda, Drumstick, and Mud turtle. A red X indicates failed classification.

Classification + Localization:


- Steel drum:** A video frame showing a person playing a steel drum. Bounding boxes and labels are shown: Version cat (red), Picket fence (cyan), Steel drum (green), Folding chair (orange), and Loud speaker (blue). A green checkmark indicates successful classification and localization.

ConvNet Configuration					
A	A-1	B	D	E	F
11 weight layers	13 weight layers	16 weight layers	16 weight layers	17 weight layers	17 weight layers
conv-5-64	conv-5-64	conv-5-64	conv-5-64	conv-5-64	conv-5-64
maxpool	maxpool	conv-5-128	conv-5-128	conv-5-128	conv-5-128
conv-5-128	conv-5-128	conv-5-128	conv-5-128	conv-5-128	conv-5-128
		maxpool	conv-5-128	conv-5-128	conv-5-128
conv-3-256	conv-3-256	conv-3-256	conv-3-256	conv-3-256	conv-3-256
conv-3-256	conv-3-256	conv-3-256	conv-3-256	conv-3-256	conv-3-256
		maxpool			
conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512
conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512
		maxpool			
conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512
conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512	conv-5-512
		maxpool			
FC-4096	FC-4096	FC-4096	FC-4096	FC-4096	FC-4096
FC-1000	FC-1000	FC-1000	FC-1000	FC-1000	FC-1000
soft-max	soft-max	soft-max	soft-max	soft-max	soft-max




More Applications

Applications – Object Detection



Ren S, He K, Girshick R, et al. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. NIPS 2015.

Applications – Object Detection



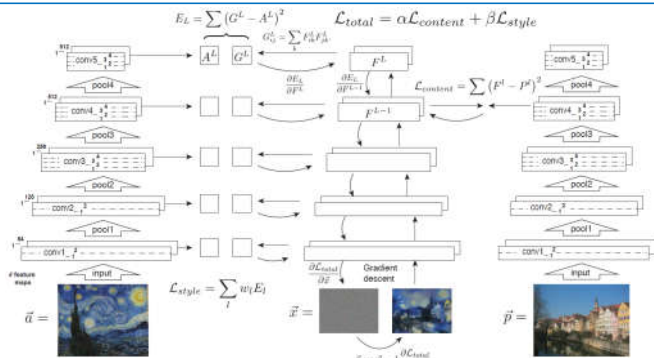
Ren S, He K, Girshick R, et al. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. NIPS 2015.

Applications – Style Transfer



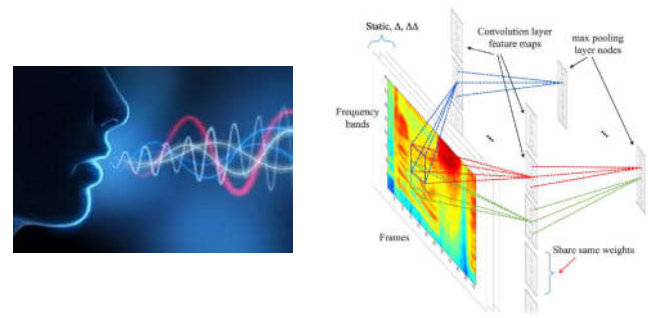
Gatys L A, Ecker A S, Bethge M. Image Style Transfer Using Convolutional Neural Networks. CVPR, 2016.

Applications – Style Transfer




Gatys L A, Ecker A S, Bethge M. Image Style Transfer Using Convolutional Neural Networks. CVPR, 2016.

Applications – Speech Recognition



Abdel-Hamid O, Mohamed A R, Jiang H, et al. Convolutional Neural Networks for Speech Recognition. IEEE/ACM Transactions on Audio Speech & Language Processing, 2014, 22(10):1533-1545.

Applications – Text Classification



Yoon Kim. Convolutional Neural Networks for Sentence Classification. EMNLP, 2014.

Materials

- Paper: Gradient-based Learning Applied to Document Recognition
Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner, 1998.
- Paper: Deep Learning
Y. Lecun, Y. Bengio, and G. Hinton, *Nature*, 2015.
- Course: Convolutional Neural Networks for Visual Recognition
<http://cs231n.stanford.edu/>
- Tool: CNN Visualization
<https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html>

Questions

